

Vamshi Krishna Bonagiri

✉ vamshi.12.2003@gmail.com

🌐 [Homepage](#)

🌐 [LinkedIn](#)

🔍 [Google Scholar](#)

PhD student working on representation geometry and interpretability for understanding and aligning LLMs, with prior work spanning AI safety, evaluation, and ethics.

Education

- 2025–present **PhD in Natural Language Processing**
Mohammed Bin Zayed University of Artificial Intelligence (MBZUAI), Abu Dhabi
- 2020–2025 **B.Tech. in Computer Science and M.S. by Research in Computational Linguistics,**
International Institute of Information Technology, Hyderabad GPA: 9.01/10

Research Experience

- 2025 **Center for Human-Compatible AI (CHAI), UC Berkeley, Research Intern, USA**
Worked at [Stuart Russell's](#) Lab under the supervision of [Benjamin Plaut](#) on evaluating and improving LLM Agent uncertainty quantification methods for complex multi-turn settings.
- 2022–2024 **Microsoft Research, Visiting Researcher, India**
Collaborated with [Monojit Choudhury](#) on creating a benchmark on code-mixed acceptability and with [Tanuja Ganu](#) on developing cross-modal adversarial attacks.
- 2021–2025 **Precog, IIIT Hyderabad, Student/Research Assistant, India**
Conducted research on NLP, AI safety, and trustworthy ML under the supervision of [Ponnu-rangam Kumaraguru](#).
- 2023–2024 **KAI² Lab, University of Maryland, Undergraduate Research Assistant, USA**
Worked on LLM Evaluations for reliability and trustworthiness, led by [Manas Gaur](#).
- 2021–2022 **Tokyo Institute of Technology, Research Intern, Japan**
Analyzed a corpus of 87,000 Reddit comments related to Deepfakes with [Sasahara](#).

Industry Experience

- 2024 **Sprinklr, Machine Learning Product Engineering Intern, Gurgaon, India**
Developed a custom LLM-based Multi-Agent framework (similar to [Autogen](#)) with **30% higher accuracy, 85% latency reduction, and 80% cost reduction** compared to off-the-shelf solutions, **accelerating sales analysis by 28x**.
- 2022 **Observe.ai, Machine Learning Research Intern, Bangalore, India**
Developed a low-resource multilingual text classification architecture with a **75+ F1 score, enabling product adaptation** from English to Spanish.
- 2022 **Smart City Living Labs, Software Engineer Intern, Hyderabad, India**
Developed and [deployed an immersive 3D dashboard](#) for real-time monitoring of 300+ IoT devices across the IIIT Hyderabad campus.

Teaching & Mentoring

- 2026 **SPAR, Mentor**
Currently supervising [projects on agentic situational awareness and efficient agent safety evals](#)

- 2024-2025 **IIIT Hyderabad & NPTEL**, *Teaching Assistant*
 (1)“Using AI Tools for Research” Workshop at IIT Hyderabad (2025) (2) Responsible & Safe AI NPTEL course (2024) (3) Responsible and Safe AI Systems course, supported by Open Philanthropy grant (2024) (4) Introduction to NLP (2023)
- 2024 **ACM India Summer School on Responsible & Safe AI**, *Teaching Assistant*, IIT Madras
- 2024 **Bluedot Impact**, *AI Safety Fundamentals Facilitator*, United Kingdom
 Facilitated an international cohort in the [AI Safety fundamentals course](#).

Grants & Awards

- 2023-2025 **Research Distinctions:** OpenAI Research Access Grant, Google DeepMind Symposium 2025 (top 200 ML scholars in India), [IndoML 2024](#) (top 20 ML graduate scholars in India), Open Philanthropy AI Safety Course Grant, [AI Alignment Workshop Grant by FAR.AI](#), [Black Box NLP Grant](#), [Global Challenges Project](#), Adaption Labs Research Grant
- 2020-2024 **Academic Awards:** Dean’s List (Spring 2021, 2022, 2024 & Monsoon 2020, 2021), Research Award (Spring 2024)
- 2017-2021 **Competition Awards:** Megathon 2021 Winner (COVID mask detector), NASA AMES Space Settlement Contest 2017 (Top 3 internationally)

Leadership & Service

- 2025 **Overall Co-ordinator**, Reading Group, NLP Dept., MBZUAI
- 2021-2023 **Overall Co-ordinator**, Open Source Developers Group, IIIT Hyderabad
- 2020-2022 **Tech Team Member**, Entrepreneurship Cell, IIIT Hyderabad
- 2024-2025 **Volunteer**, Mental Health Support Group

Publications

- 2026 ***Intrinsic Guardrails: How Semantic Geometry of Personality Interacts with Emergent Misalignment in LLMs***
 Aneja, K., Mittal, M., Goel, A., Kumaraguru, P., **Bonagiri, V.K.**
AI4GOOD, ICML 2026
- 2026 ***Efficient Safety Benchmarking via Item Response Theory***
 Spagliardi, F.*, Silva, M.*, Datta, A.*, Zhou, A., **Bonagiri, V.K.**, Cruz, D.
AI4GOOD, ICML 2026
- 2026 ***If Pigs Could Fly... Can LLMs Logically Reason Through Counterfactuals?***
 Joishy, A.R.*, Balappanawar, I.B.*, **Bonagiri, V.K.***, Gaur, M., Thirunarayan, K., Kumaraguru, P.
IEEE WCCI, IJCNN 2026
- 2025 ***Check Yourself Before You Wreck Yourself: Selectively Quitting Improves LLM Agent Safety***
Bonagiri, V.K., Kumaraguru, P., Nguyen, K., Plaut, B.
Regulatable ML and Reliable ML Workshops, NeurIPS 2025
- 2025 ***Dark Side of the Tune: Investigating the maladaptive outcomes of excessive music consumption in the age of unlimited music access***

Bonagiri, V.K., Alluri, V.

18th International Conference on Music Perception and Cognition (ICMPC) 2025

- 2024 *Evaluating Moral Consistency in Large Language Models*
Bonagiri, V.K., Vennam, S., Govil, P., Kumaraguru, P., and Garg, M.
LREC-COLING 20-25 May, 2024, Torino, Italy
- 2025 *From Human Judgements to Predictive Models: Unravelling Acceptability in Code-Mixed Sentences*
Kodali, P., Goel, A., **Bonagiri, V.K.**, Choudhury, M., Shrivastava, M., Kumaraguru, P.
ACM TALLIP (Journal) 2025
- 2025 *COBIAS: Contextual Reliability in Bias Assessment*
Govil, P., **Bonagiri, V.K.**, Garg, M., and Kumaraguru, P
ACM Web Science 2025
- 2024 *K-PERM: Personalized Response Generation Using Dynamic Knowledge Retrieval and Persona-Adaptive Queries*
Raj, K., Roy, K., **Bonagiri, V.K.**, Thirunarayanan, K
AAAI-MAKE 2024
- 2023 *Representation Learning for Identifying Depression Causes in Social Media*
Govil, P., **Bonagiri, V.K.**, Garg, M., and Kumaraguru, P
Proceedings of KDD KiL 2023. Long Beach, California, USA August 6, 2023
- 2023 *Towards Effective Paraphrasing for Information Disguise. In European Conference on Information Retrieval (pp. 331-340)*
Agarwal, A., Gupta, S., **Bonagiri, V.K.**, Gaur, M., Reagle, J., Kumaraguru, P
ECIR March, 2023
- 2022 *Are deepfakes concerning? analyzing conversations of deepfakes on Reddit and exploring societal implications. In Proceedings of the 2022 CHI Conference on Human Factors in Computing Systems (pp. 1-19)*
Gamage, D., Ghasiya, P., **Bonagiri, V.K.**, Whiting, M. E., Sasahara, K
CHI April, 2022