Vamshi Krishna Bonagiri

vamshi.12.2003@gmail.com Homepage LinkedIn Google Scholar

Education

2020–2025 **B.Tech. in Computer Science and M.S. by Research in Computational Linguistics,**
International Institute of Information Technology, Hyderabad GPA: 9.10/10

Research Experience

2025-Present **Center for Human-Compatible AI (CHAI), UC Berkeley**, *Visiting Research Scholar*, USA
Working with Dr. Stuart Russell and Benjamin Plaut on evaluating and improving LLM Agent uncertainty quantification methods for complex multi-turn settings

2022-2024 **Microsoft Research**, *Visiting Researcher*, India
Collaborated with Dr. Monojit Choudhury on creating a benchmark on code-mixed acceptability and with Tanuja Ganu on developing cross-modal adversarial attacks.

2021-2025 **Precog, IIIT Hyderabad**, *Research Assistant/Undergraduate Researcher*, India
Worked on NLP and AI Safety Research under the supervision of Prof. Ponnurangam Kumaraguru, Here's a list of published research.

2023–2024 **KAI^2 Lab, University of Maryland**, *Undergraduate Research Assistant*, USA
Worked on LLM Evaluations for reliability and trustworthiness, led by Dr. Manas Gaur.

2021-2022 **Tokyo Institute of Technology**, *Research Intern*, Japan
Analyzed a corpus of 87,000 Reddit comments related to Deepfakes with Dr.Sasahara.

Industry Experience

2024 **Sprinklr**, *Machine Learning Product Engineering Intern*, Gurgaon, India
Developed a custom LLM-based Multi-Agent framework (similar to Autogen) with **30% higher accuracy**, **85% latency reduction**, and **80% cost reduction** compared to off-the-shelf solutions, **accelerating sales analysis by 28x**.

2022 **Observe.ai**, *Machine Learning Research Intern*, Bangalore, India
Developed a low-resource multilingual text classification architecture with a **75+ F1 score**, **enabling product adaptation** from English to Spanish.

2022 **Smart City Living Labs**, *Software Engineer Intern*, Hyderabad, India
Developed and deployed an immersive 3D dashboard for real-time monitoring of 300+ IoT devices across the IIIT Hyderabad campus.

Teaching & Mentoring

2024-2025 **IIIT Hyderabad & NPTEL**, *Teaching Assistant*
(1) Using AI Tools for Research Workshop at IIT Hyderabad (2025) (2) Responsible & Safe AI NPTEL course (2024) (3) Responsible and Safe AI Systems course, supported by Open Philanthropy grant (2024) (4) Introduction to NLP (2023)

2024 **ACM India Summer School on Responsible & Safe AI**, *Teaching Assistant*, IIT Madras

2024 **Bluedot Impact**, *AI Safety Fundamentals Facilitator*, United Kingdom
Led an international cohort in the AI Safety fundamentals course.

last updated: May 24, 2025

Grants & Awards

2023-2025 **Research Distinctions**: Google DeepMind Symposium 2025 (top 200 ML scholars in India), IndoML 2024 (top 20 ML graduate scholars), Research Award, IIIT Hyd (Spring 2024), Open Philanthropy AI Safety Course Grant, AI Alignment Workshop Grant by FAR.AI, Black Box NLP Grant, Global Challenges Project

2020-2024 **Academic Awards**: Dean's List 2 (Spring 2024, 2022), Dean's List 1 (Monsoon 2021), Merit List (Spring 2021, Monsoon 2020)

2017-2021 **Competition Awards**: Megathon 2021 Winner (COVID mask detector), NASA AMES Space Settlement Contest 2017 (Top 3 internationally)

Leadership & Service

2021-2023 **Overall Co-ordinator**, Open Source Developers Group, IIIT Hyderabad

2020-2022 **Tech Team Member**, Entrepreneurship Cell, IIIT Hyderabad

2024-2025 **Volunteer**, Mental Health Support Group

2020-2025 **Freelancer** - Projects in blog writing, book reviews, poetry generation, web development, and Machine Learning tutoring

Papers and Publications

2025 ***If Pigs Could Fly... Can LLMs Logically Reason Through Counterfactuals?***
Bonagiri, V.K. et al.
Under Review at Neurips *2025*

2025 ***From Human Judgements to Predictive Models: Unravelling Acceptability in Code-Mixed Sentences***
Kodali, P., Goel, A., Bonagiri, V. K., Choudhury, M., Shrivastava, M., Kumaraguru, P.
ACM TALLIP (Journal) *2025*

2025 ***COBIAS: Contextual Reliability in Bias Assessment***
Govil, P., Bonagiri, V., Garg, M., and Kumaraguru, P
ACM Web Science 2025

2024 ***Evaluating Moral Consistency in Large Language Models***
Bonagiri, V., Vennam, S., Govil, P., Kumaraguru, P., and Garg, M. SaGE
LREC-COLING *20-25 May, 2024, Torino, Italy*

2024 ***K-PERM: Personalized Response Generation Using Dynamic Knowledge Retrieval and Persona-Adaptive Queries***
Raj, K., Roy, K., Bonagiri, V., Thirunarayanan, K
AAAI-MAKE *2024*

2023 ***Representation Learning for Identifying Depression Causes in Social Media***
Govil, P., Bonagiri, V., Garg, M., and Kumaraguru, P
Proceedings of KDD KiL 2023. Long Beach, California, USA *August 6, 2023*

2023 ***Towards Effective Paraphrasing for Information Disguise. In European Conference on Information Retrieval (pp. 331-340)***
Agarwal, A., Gupta, S., Bonagiri, V., Gaur, M., Reagle, J., Kumaraguru, P
ECIR *March, 2023*

2022

last updated: May 24, 2025

Are deepfakes concerning? analyzing conversations of deepfakes on Reddit and exploring societal implications. In Proceedings of the 2022 CHI Conference on Human Factors in Computing Systems (pp. 1-19)
Gamage, D., Ghasiya, P., Bonagiri, V., Whiting, M. E., Sasahara, K
CHI *April, 2022*